Do you need your own OAuth app?

A lookup tool that tells a Composio developer, in about thirty seconds, whether they need their own OAuth app, how long approval really takes, and what they can start today.

Live	composio-2.vercel.app
Repo	github.com/chayan171999/composio-2-

What it is

Three hand-verified seed guides — GitHub, Google/Gmail and Meta/Facebook — chosen to cover three different blocker types: no review, scope review, and business verification. Anything else is researched live by an agent that reads the provider's own documentation and writes a guide in the same schema, streaming its progress. Slack, Zoom and Figma are agent-generated examples. A new guide takes 10–15 seconds.

What's worth a closer look

Evidence is enforced in code, not prompting. Every factual statement is a claim with a status (tested / sourced / inferred / unverified), its sources, and the date checked. The agent fetches documentation pages itself, so the citation allow-list is built from HTTP statuses the server actually saw — a page that 404s cannot be cited, whatever the model writes. Generated guides are forbidden from claiming *tested*, and dangling claim references are rejected; anything failing goes back for one repair pass.

The Composio managed-OAuth answer is read from data. Composio's public toolkit page embeds the same field its authenticated API returns, so that question is answered deterministically rather than by a model reading an excerpt — which got it backwards, because the page is over 500k characters and the excerpt landed on an unrelated FAQ entry.

Two things to know before you try it

- **It runs on Groq, not the Anthropic API the brief specified.** Anthropic has no free tier and "no paid accounts" was the harder constraint. Groq's free tier allows roughly 1,000 model requests a day — two or three per guide — but meters 8,000 tokens per minute, so the practical pace is about one generation a minute. If several are run back to back the agent waits for the limit to clear and says so, rather than failing.
- **That 8,000-token ceiling applies per request, including the output**, so at most four documentation pages reach the model. It is the main quality constraint and it has a real cost: on one Zoom run the agent reported no review where the honest answer is scope review, having not retrieved the Marketplace pages. **A generated guide can miss a blocker.**
- **Generated guides currently have no shareable URL.** No key-value store is configured on this deployment, so a new guide is rendered on the page rather than saved. The six guides already in the library are files in the repo and always load.

Where the detail lives

`README.md` covers setup, the decisions behind the guide's ordering and evidence model, how the agent works, and its limits — including what is not built. `docs/verification.md` records what was checked against primary sources, the corrections that turned up, and what remains unverified.

Next.js (App Router) · TypeScript · Tailwind · Zod · Groq · Vercel. To run locally: `npm install`, copy `.env.example` to `.env.local` and set `GROQ_API_KEY` (free, no billing account), then `npm run dev`.